

Reinforcement Learning

Dynamic Programming Sutton/Barto* Chapter 4

Michael Hahsler, SMU

With figures from Sutton/Barto*

*Sutton and Barto, Reinforcement Learning: An Introduction,
2nd edition, MIT Press, Cambridge, MA, 2018



Summary of Notation

General

X	capital letters: random variables
x, p	lower-case letters: realizations of random variables or scalar functions
$\mathbf{w}$	Bold lower-case letters: real-valued vectors (even if random variables)
$\mathbf{W}$	bold capitals: matrices
α	Greek letters: parameters (vectors if in bold)
$\Pr\{X = x\}$	probability that a random variable X takes on the value x
$X \sim p$	random variable X selected from distribution $p(x) = \Pr\{X = x\}$
$\mathbb{E}[X]$	expectation of a random variable X , i.e., $\mathbb{E}[X] = \sum_x p(x)x$
$\operatorname{argmax}_a f(a)$	a value of action a at which $f(a)$ takes its maximal value

Value Function

G_t	return (cumulative reward) following time t
$G_{t:h}$	return from t to h (discounted and corrected)
$v_\pi(s)$	value of state s under policy π (expected return)
$v_*(s)$	value of state s under the optimal policy
$q_\pi(s, a)$	value of taking action a in state s under policy π
$q_*(s, a)$	value of taking action a in state s under the optimal policy
V, V_t	array estimates of state-value function v_π or v_*
Q, Q_t	array estimates of action-value function q_π or q_*

MDP

s, s'	states
a	an action
r	a reward
$\mathcal{S}$	set of all (nonterminal) states, $\mathcal{S}^+$ are all states
$\mathcal{A}(s)$	set of all actions available in state s
γ	discount-rate parameter
t	discrete time step
T	final time step of an episode (a.k.a. horizon)
A_t	random variable for the action at time t
S_t	random variable for the state at time t
R_t	random variable for the reward at time t
$p(s', r s, a)$	probability of transition to state s' and receiving reward r , from state s taking action a .
$p(s' s, a)$	probability of transition to state s' from state s taking action a .
$r(s, a)$	expected immediate reward from state s after action a .
$r(s, a, s')$	expected immediate reward from state s to s' with action a .
$\pi(a s)$	probability of taking action a in state s under stochastic policy π
$\pi(s)$	action taken in state s under deterministic policy π

Goal

- **Remember:** MDPs have an optimal deterministic policy.
- **Goal:** Compute the **optimal deterministic policy** $\pi_* = \operatorname{argmax}_a q_*(s, a)$ given a perfect finite MDP model.
- Solve the Bellman optimality equations (the optimal policy always picks the best action):

$$\begin{aligned} v_*(s) &= \max_a \mathbb{E}_\pi [R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a] \\ &\text{or} \\ q_*(s, a) &= \max_a \mathbb{E}_\pi \left[R_{t+1} + \gamma \underbrace{\max_{a'} q_*(S_{t+1}, a')}_{= v_*(S_{t+1})} \middle| S_t = s, A_t = a \right] \end{aligned}$$

Control: Find the optimal policy

- **Issues:** The value estimate for S_t depends on the current estimates of the successor states S_{t+1} ! We need to solve all optimality equations at the same time.

- **Idea:** We can find the optimal policy by stepwise **improving a policy** till it is optimal. A policy is better if it has larger state values.

$$\pi \geq \pi' \Leftrightarrow v_\pi(s) \geq v_{\pi'}(s) \text{ for all } s \in \mathcal{S}$$

- To do this we need to be able to **evaluate a policy** (i.e., compute v_π and $v_{\pi'}$).

Prediction/Evaluation:
Estimate the value function for a policy.

Tabular Methods

- Tabular means: The value functions v or q are stored in arrays (tables) called V or Q .
- Important property: Entries in the table do not influence each other. We can change them independently.
- Tabular methods find the exact solution (i.e., optimal value function and policy) given enough memory and time.
- **Issues:**
 - Table size:
 - v has $|\mathcal{S}|$ entries
 - q has $|\mathcal{S}| \times |\mathcal{A}|$ entries

This is only feasible for problems with a small discrete state/action space!
 - No generalization: Knowledge about the value of one state (state-action pair) does not help us with similar states
- We will later talk about approximate methods that generalize and can work with large state spaces.

Value Function

s	State Value $V(s)$
(1,1)	0.7453
(1,2)	0.8016
...	...

Q-Table

s	a	$q(s, a)$
(1,1)	up	0.7453
(1,1)	right	0.6709
(1,1)	down	0.7003
(1,1)	left	0.7109
...	...	...

or

	up	right	down	left
(1,1)	0.7453	0.6709	0.7003	0.7109
...	...	...	...	...
...	...	...	...	...

Prediction vs. Control

- In this class, we will always talk about two tasks:
 1. **Prediction to evaluate a policy:** Estimate the value function for a given policy. We can use this to compare two policies.
 2. **Control:** Control the agent's behavior to find the optimal policy. This requires being able to evaluate policies.

Prediction: Policy Evaluation

What is the value function of a given policy?

Policy Evaluation to Estimate v_π

- **Goal:** Compute the state-value function v_π for a given policy π .
- Bellman expectation equations give us the recursive relationship:

$$\begin{aligned} v_\pi(s) &\stackrel{\text{def}}{=} \mathbb{E}_\pi[G_t | S_t = s] \\ &= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= \mathbb{E}_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')] \end{aligned}$$

with $|S|$ equations, one for each state s .

- **Guarantee:** There exists a unique v_π if $\gamma < 1$ or termination is guaranteed by π .

Solution Option 1: Linear Program

- Directly solve the system of $|\mathcal{S}|$ linear equations.

$$\begin{aligned}v_{\pi}(s_1) &= \sum_a \pi(a|s_1) \sum_{s',r} p(s',r | s_1, a) [r + \gamma v_{\pi}(s')] \\v_{\pi}(s_2) &= \sum_a \pi(a|s_2) \sum_{s',r} p(s',r | s_2, a) [r + \gamma v_{\pi}(s')] \\&\dots \text{ for all } s \in \mathcal{S}\end{aligned}$$


- This is hard if $|\mathcal{S}|$ is very large.

Solution Option 2: Iterative Policy Evaluation

Dynamic Programming: solves a complex optimization problem by breaking it down into smaller, overlapping subproblems and using their solutions to build the overall optimal solution.

- Start with arbitrary values for all $v(s)$
- Update all state values (called sweeping over the state space) using the **Bellman Update** (k is the iteration counter):

Backup diagram




$$v_{k+1}(s) \stackrel{\text{def}}{=} [R_{t+1} + \gamma v_k(S_{t+1}) | S_t = s]$$
$$= \sum_a \pi(a|s) \underbrace{\sum_{s',r} p(s',r | s, a) [r + \gamma v_k(s')]}_{\text{In DP called "expected update."}}$$


In DP called "expected update."
Uses the expectation over all possible actions and next states.

- **Convergence:** Updates converge to v_π for $k \rightarrow \infty$ iterations.

Sweep the state space



Bellman updates



Convergence of Iterative Policy Evaluation

- The Bellman update can be rewritten as an operator:
Bellman Expectation Operator T^π :

$$(T^\pi v)(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r | s, a) [r + \gamma v(s')]$$

- The operator T^π is γ -contracting with v_π as its fixed point:
 $\|T^\pi v_\pi - T^\pi v\|_\infty \leq \gamma \|v_\pi - v\|_\infty$ for any $\gamma < 1$

Notation:

$(T^\pi v)$ is the new value function after applying the operator. We could write $v_{k+1} = T^\pi(v_k)$

Note:

$$v_\pi = T^\pi(v_\pi)$$

This means the largest difference (infinity norm) will be reduced, leading to convergence of $v = v_\pi$ when T^π is applied many times.

Convergence is only guaranteed if the discount factor is strictly smaller than 1!

Alg. Iterative Policy Evaluation

Notation: V is the tabular estimate of v

Iterative Policy Evaluation, for estimating $V \approx v_\pi$

Input π , the policy to be evaluated

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$ arbitrarily, for $s \in \mathcal{S}$, and $V(\text{terminal})$ to 0

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$

Bellman operator T^π :
 $V_{k+1} = T^\pi V_k$

Stop when value function changes little:

$$\|T^\pi V_k - V_k\|_\infty < \theta$$

It can be shown that the error is bounded:

$$\|V_k - v_\pi\|_\infty < \frac{\theta}{1 - \gamma}$$

If we stop policy evaluation early (after n iterations) then it is called **modified policy iteration** and we end up with an approximation.

Policy Iteration

Find the optimal policy based on the environment model.

Policy Improvement

- **Goal:** Find a better policy.

- **Policy improvement theorem:** If

$$q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s) \quad \forall s \in \mathcal{S}$$

then π' is at least as good or better than π , i.e.,

$$v_{\pi'}(s) \geq v_{\pi}(s) \quad \forall s \in \mathcal{S}.$$

- **Policy improvement:** create a new policy π' by choosing in each state the best (aka greedy) action with respect to $q_{\pi}(s, a)$:

$$\pi'(s) = \underset{a}{\operatorname{argmax}} q_{\pi}(s, a)$$

- **Guarantee:** This **greedy policy** π' is strictly better unless π is already optimal.

Policy Iteration

Idea: Create a sequence of monotonically improving policies:

$$\pi_0 \xrightarrow{E} v_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} v_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} v_{\pi_2} \xrightarrow{I} \pi_3 \xrightarrow{E} v_{\pi_3} \xrightarrow{I} \dots \pi_* \xrightarrow{E} v_*$$

$\xrightarrow{E}$ policy eval. step
 $\xrightarrow{I}$ policy improvement step

Guarantees:


- Policy improvement is guaranteed to improve the policy.
- Finite MDPs have a finite number of policies $\rightarrow$ Policy Iteration must **converge**!

Advantages:

- Policy becomes stable before v_* is reached.
- Policy iteration typically converges fast.

Issue:

- Policy estimation is computationally expensive!
Using fewer iterations leads to modified policy iteration.



Alg. Policy Iteration

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$; $V(\text{terminal}) \doteq 0$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)

Just policy evaluation from before.
We can also stop this loop early
(modified policy iteration).

3. Policy Improvement

policy-stable $\leftarrow$ *true*

For each $s \in \mathcal{S}$:

old-action $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

If *old-action* $\neq \pi(s)$, then *policy-stable* $\leftarrow$ *false*

If *policy-stable*, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

Picks the greedy action with the highest
Q-value to improve the policy.

Value Iteration

Find the optimal value function based on the environment model.

Value Iteration

- Combines policy improvement and truncated policy evaluation into a single step. This is called the **expected Bellman optimality update**.


$$\begin{aligned} v_{k+1}(s) &\stackrel{\text{def}}{=} \max_a \mathbb{E}[R_{t+1} + \gamma v_k(S_{t+1}) | S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_k(s')] \end{aligned}$$

$\max_a$... policy improvement
 $\mathbb{E}$... truncated policy evaluation (1 step)

- Written as the **Bellman optimality operator**:

$$(T^*v)(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v(s')]$$

- Guarantee:** T^* is a non-linear operator, but it is also γ -contracting. The Sequence $\{v_k\}$ converges to v_* .



Alg. Value Iteration

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$; $V(\text{terminal}) \doteq 0$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$

3. Policy Improvement

For each $s \in \mathcal{S}$:

$\text{old-action} \leftarrow \pi(s)$

$\pi(s) \leftarrow \arg\max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

If $\text{old-action} \neq \pi(s)$, then $\text{policy-stable} \leftarrow \text{false}$

If policy-stable , then stop and return $V \approx v_*$

Combine improvement and modified evaluation (1 iteration) into a single update.

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$

Output a deterministic policy, $\pi \approx \pi_*$, such that

$\pi(s) = \arg\max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

Bellman optimality operator:

$$V_{k+1} = T^*V_k$$

Stopping criterion from policy evaluation

Extract the greedy policy.

Issues with DP Methods

Requires perfect knowledge of the model!

Space Complexity

- Tabular method: stores the complete value function as table V .

Time Complexity

- Updates repeatedly sweep over the whole state space.
- Wastes computation by updating states that are not relevant for the optimal behavior. I.e., the optimal policy will never visit them.


GPI: Generalized Policy Iteration

Generalizing value and policy iteration into a single framework.

GPI: Generalized Policy Iteration

Steps


- **Policy evaluation:** make the value function consistent with the current policy π .
- **Policy improvement:** make policy greedy with respect to the current value function estimate V .



Most DP methods can be described as a special case of GPI:

- **Policy iteration:** each evaluation and improvement step is completed (by a sweep over $\mathcal{S}$) before moving on.
- **Value iteration:** uses a single iteration for the policy evaluation step.
- **Asynchronous DP:** Instead of sweeps over $\mathcal{S}$, updates states in some other order (e.g., focus on states where the value function changed a lot in the last step). This interleaves pol. eval. and pol. Improvement.

All three approaches typically converge to the optimal solution.



Efficiency

- **Direct Search:** Search space $O(|\mathcal{A}|^{|\mathcal{S}|})$ is **exponential** in the number of states $|\mathcal{S}|$.
- **LP (linear program):** Solve a system of $|\mathcal{S}|$ equations with $|\mathcal{S}|$ state variables and $|\mathcal{S}| \times |\mathcal{A}|$ constraints has a time complexity of $O(|\mathcal{S}|^3)$.
- **DP:** Time $O(|\mathcal{S}|^2 \times |\mathcal{A}|)$ is **quadratic** in $|\mathcal{S}|$.

Search and LP are impractical. The number of actions is typically relatively small. **DP can be used today with millions of states.**

What you need to know



Requirements:

- **Model access:** DP methods need access to the **model's** function p .
- **Small discrete state/action space** since it is a tabular method.

Tasks:

- **Evaluation/Prediction:** Calculates the value function for a given policy.
- **Control:** Calculates the optimal policy for an MDP.

Approach:

- Convert the recursive Bellman equations into the Bellman update operators and sweep the state space till the optimal solution is found.
 - Expected Bellman update operator T^π vs. Bellman optimality operator T^*
 - Value iteration is a special case of modified policy iteration with a single policy evaluation iteration.
 - Convergence requires either $\gamma < 1$ or a guarantee that the policy will reach a terminal state.

GPI:

- All DP methods can be generalized using the GPI evaluation/improvement framework.